nonprobsvy – An R package for modern methods for non-probability surveys

Łukasz Chrostowski
Analyx

Piotr Chlebicki ⓘ
Stockholm University

Maciej Beręsewicz ⓘ
Poznań University of Economics and Business
Statistical Office in Poznań

Abstract

The following paper presents **nonprobsvy** – an R package for inference based on non-probability samples. The package implements various approaches that can be categorized into three groups: prediction-based approach, inverse probability weighting and doubly robust approach. Contrary to the existing packages **nonprobsvy** assumes existence of either full population or probability-based population information and laverage the **survey** package for inference. The package implements both analytical and bootstrap variance estimation for all of the proposed estimators. In the paper we present the theory behind the package, its functionalities and case study that showcases the usage of the package. The package is aimed at official staticans, public opinion or market researchers who whould like to use non-probability samples (e.g. big data, opt-in web panels, social media) to accurately estimate popualtion characteristics.

Keywords: data integration, doubly robust estimation, propensity score estimation, mass imputation, **survey**.

1. Introduction

In official statistics, information about the target population and its characteristics is mainly collected through probability surveys, censuses or is obtained from administrative registers and covers all (or nearly all) units of the population. However, owing to increasing non-response rates, particularly unit non-response and non-contact, which result from the growing respondent burden as well as rising costs of surveys conducted by National Statistical Insti-

tutes, non-probability data sources are becoming more popular (Beręsewicz 2017; Beaumont 2020; Biffignandi and Bethlehem 2021). Non-probability surveys, such as opt-in web panels, social media, scanner data, mobile phone data or voluntary register data, are currently being explored for use in the production of official statistics (Citro 2014; Daas, Puts, Buelens, and Hurk 2015), public opinion studies (Schonlau and Couper 2017) or market research (cf. Grow, Perrotta, Del Fava, Cimentada, Rampazzo, Gil-Clavel, Zagheni, Flores, Ventura, and Weber 2022). Since the selection mechanism underlying these sources is unknown, standard design-based inference methods cannot be directly applied and, in the case of large datasets, can lead to the *big data paradox* described by Meng (2018).

Table 1 compares basic characteristics of probability and non-probability samples. In particular, it shows the advantages and disadvantages of each type with respect to the selection mechanism, the population coverage, bias, variance, costs and timeliness. In general, the quality of non-probability samples suffers from an unknown selection mechanism (i.e. unknown probabilities of inclusion) and under-coverage of certain groups from the population (e.g. older people). As a result, direct estimates based on non-probability samples are biased and, in most cases, are characterised by small variance owing to their size, which is known as the *big data paradox*, i.e. the larger the sample, the larger the bias. Certainly, the costs and timeliness of these surveys are significantly smaller than those of probability samples.

Factor	Probability sample	Non-probability sample
Selection	Known probabilities	Unknown self-selection
Coverage	Complete	May be incomplete
Estimation bias	Unbiased under design	Potential systematic bias
Variance of estimates	Typically high	Typically low
Cost	High	Low
Timeliness	Long delay	Very short delay

Table 1: A comparison of probability and non-probability samples and their characteristics

To address this problem, several approaches have been proposed, which rely on the estimation of propensity scores (i.e. inclusion probabilities) for deriving inverse probability weights (IPW; also known as propensity score weighting/adjustment, cf. Lee (2006); Lee and Valliant (2009)), on model-based prediction (in particular, mass imputation estimators; MI) and on the doubly robust (DR) approach involving IPW and MI estimators. Two main scenarios are usually considered: 1) only population-level means or totals are available, and 2) unit-level data are available either in the form of registers covering the whole population or in the form of probability surveys (cf. Elliott and Valliant 2017). Wu (2022) classified these approaches into three groups that require a joint randomization framework involving a *probability sampling design* (denoted as p) and an outcome regression model (denoted as ξ) or a propensity score model (denoted as q). According to this classification, IPW estimators represent the qp framework, MI estimators represent the ξp framework, and DR estimators can represent either the qp or the ξp framework.

Most approaches assume that population data is available in order to reduce the bias of non-probability sampling by reweighting to reproduce known population totals/means (i.e. IPW estimators); by modelling the target variable using various techniques (i.e. MI estimators); or by combining both approaches (e.g. DR estimators, cf. Chen, Li, and Wu (2020); see also Multilevel Regression and Post-stratification, MRP; *Mister-P*, cf. Gelman (1997)). This topic

has become very popular and a number of new methods have been proposed; for instance non-parametric approaches based on nearest neighbours (Yang, Kim, and Hwang 2021), kernel density estimation (Chen, Yang, and Kim 2022), empirical likelihood (Kim and Morikawa 2023), model-calibration with LASSO (Chen, Valliant, and Elliott 2018) or quantile balanced IPW (Beręsewicz, Szymkowiak, and Chlebicki 2025) to name a few. It should be highlighted that, in contrast to probability samples, there is no single method that can be used for non-probability samples. Based on the methods available in the literature several statistical software solutions have been developed. These pieces of software will be presented and compared with **nonprobsvy** in the next section.

1.1. Software for non-probability samples

Table 2 presents a comparison of selected packages in terms of the availability of various inference methods. We focus on packages available through CRAN or PyPI (for non-CRAN or non-PyPI software see Cobo, Ferri-García, Rueda-Sánchez, and Rueda (2024)). The comparison includes four packages that particularly focus on non-probability samples: in R – **NonProbEst** (Rueda, Ferri-García, and Castro 2020) and our **nonprobsvy**, and in Python – **balance** (Sarig, Galili, and Eilat 2023), **inps** (Castro Martín 2024). In addition, we have included two R packages that implement specific methods: **rstanarm** (MRP; Goodrich, Gabry, Ali, and Brilleman (2024)) and **GJRM** (generalized sample selection models; Marra and Rodicw (2023)).

Functionalities	NonProbEst	balance	inps	rstanarm	GJRM	nonprobsvy
IPW	✓	✓	✓	–	?	✓
Calibrated IPW	–	–	–	–	–	✓
MI	✓	–	–	–	–	✓
DR	–	–	✓	–	–	✓
MRP	–	–	–	✓	–	–
Sample selection	–	–	–	–	✓	–
Variable selection	✓	✓	✓	✓	✓	✓
Analytical variance	–	–	–	–	–	✓
Bootstrap variance	✓	–	–	–	–	✓
Integration with **survey** or **samplics**	–	–	–	–	–	✓

Table 2: A comparison of packages and implemented methods

The **NonProbEst** package, while more limited than **nonprobsvy**, offers various techniques, such as IPW or prediction approaches (e.g. model-calibrated). Users can choose several different settings for IPW weights, variable selection and can estimate variance using the leave-one-out Jackknife procedure. Unfortunately, the package is no longer developed (it was last updated in 2022) and some of the techniques are either outdated or have been shown to be inappropriate for non-probability samples. While the package contains functions designed for specific methods, it does not allow users to leverage the **survey** package for estimation. The **balance** package is solely dedicated to the PS approach. It assumes that the the population totals are known and the authors have implemented the variance estimator of the weighted mean as a measure of uncertainty for the IPW estimator. The weights for the IPW estimator are constructed using the approach proposed by Schonlau and Couper

(2017). The **inps** package supports the use of unit-level data from a probability sample or the population, implements IPW, MI and DR estimators, and offers users the possibility of selecting variables but is still at a very early stage of development. It also implements kernel weighting and a simple bootstrap approach via the **scipy.stats** module (Virtanen, Gommers, Oliphant, Haberland, Reddy, Cournapeau, Burovski, Peterson, Weckesser, Bright, van der Walt, Brett, Wilson, Millman, Mayorov, Nelson, Jones, Kern, Larson, Carey, Polat, Feng, Moore, VanderPlas, Laxalde, Perktold, Cimrman, Henriksen, Quintero, Harris, Archibald, Ribeiro, Pedregosa, van Mulbregt, and SciPy 1.0 Contributors 2020). Neither **balance** nor **ips** supports the use of the **samplics** module (Diallo 2021). The **GJRM** package is the only package that offers functions to estimate sample selection models used widely for correcting selection bias in observational studies (including the not-missing-at-random mechanism). Unfortunately, to best of our knowledge, there is no theory on how this approach can be used for estimating population quantities or conducting statistical inference. Finally, the MRP approach is implemented solely in the **rstanarm** package with variable selection specified by an appropriate prior.

The **nonprobsvy** package has several advantages over those presented above. Firstly, it implements state-of-the-art methods recently proposed in the literature, along with valid statistical inference procedures. Secondly, it offers other approaches, such as calibrated IPW (where PS weights match either the population or estimated totals), NN and PMM matching, various IPW and DR estimators using generalized linear models with different link functions for probability estimation. Thirdly, it supports the functions included in the **survey** package to account for the design of the probability sample. Finally, we provide a user-friendly API that mimics `glm`, `survey::svydesign` and other functions known in R, together with the main function to specify the approach and estimators. As far as we know, the **nonprobsvy** is the only software (open-access or commercial) that offers such functionalities.

The remaining part of the paper is structured as follows. Section 2 is dedicated to the theory of statistical inference based on non-probability samples. We provide the basic setup and introduce specific methods in separate subsections. We follow the notation used by Wu (2022) throughout the paper. Section 3 describes the main function and the package functionalities. Section 4 presents an empirical study showcasing the process of integrating data from the Polish Job Vacancy Survey with voluntary administrative data from the Central Job Offers Database in order to estimate the number of companies with at least one vacancy offered on a single shift. Section 5 presents classes and `S3Methods` implemented in the package. The paper ends with a summary and plans for future development. The Appendix contains Table 10 showing a list of symbols used in the paper and Section B, which presents algorithms for selected MI estimators. The Replication Materials include additional codes for specific estimators described in the paper.

2. Methods for non-probability samples

2.1. Basic setup

Let $U = \{1, ..., N\}$ denote the target population consisting of N labelled units. Each unit i has an associated vector of auxiliary variables $\boldsymbol{x}_i$ and the study (target) variable y_i. Let $\{(y_i, \boldsymbol{x}_i), i \in S_A\}$ be a non-probability sample S_A of size n_A and let $\left\{ \left(\boldsymbol{x}_i, \pi_i^B \right), i \in S_B \right\}$ be

a probability sample S_B of size n_B, where the only information known for all units in the population refers to auxiliary variables $\boldsymbol{x}$ and inclusion probabilities π^B. Each unit in the S_B sample has been assigned a design-based weight given by $d_i^B = 1/\pi_i^B$.

Let $R_i^A = I(i \in S_A)$ and $R_i^B = I(i \in S_B)$ be indicators of inclusion in the non-probability sample S_A and the probability sample S_B, respectively, which are defined for all units in the target population. Let $\pi_i^A = P(R_i^A = 1 \mid \boldsymbol{x}_i, y_i) = P(R_i^A = 1 \mid \boldsymbol{x}_i)$ be propensity scores (PS), which characterize the S_A sample's inclusion and participation mechanisms. Unlike π_i^B, the π_i^A and $d_i^A = 1/\pi_i^A$ are unknown. The description of the data is presented in a more concise form in Table 3.

Sample	ID	Inclusion (R)	Design weight (d)	Covariates ($\boldsymbol{x}$)	Study variable (y)
Non-probability	1	1	?	✓	✓
S_A	⋮	⋮	⋮	⋮	⋮
	n_A	1	?	✓	✓
Probability	1	0	✓	✓	?
S_B	⋮	⋮	⋮	⋮	⋮
	n_B	0	✓	✓	?

Table 3: Two-sample setting

The goal is to estimate the finite population mean $\mu_y = N^{-1} \sum_{i=1}^N y_i$ of the target variable y. As values of y_i are not observed in the probability sample, they cannot be used to estimate the target quantity. Instead, one could try combining the non-probability and probability samples to estimate μ_y. Given the absence of a universally accepted method for achieving this objective, assumptions vary considerably, as outlined by Wu (2022). However, the following are the main assumptions that apply to the majority of methods presented in this section:

A1 R_i^A and the study variable y_i are independent given the set of covariates $\boldsymbol{x}_i$ (i.e., $\left(R_i^A \perp y_i \right) \mid \boldsymbol{x}_i$; the MAR mechanism).

A2 All the units in the target population have non-zero PS, i.e., $\pi_i^A > 0$, $i = 1, 2, \ldots, N$ (i.e. no coverage error).

A3 The indicator variables $R_1, R_2, \ldots, R_N$ are independent given the set of auxiliary variables $(\boldsymbol{x}_1, \boldsymbol{x}_2, \ldots, \boldsymbol{x}_N)$ (i.e. no clustering).

Currently, we ignore overlap between S_A and S_B, and assume no measurement error in y_i and the fact that values of $\boldsymbol{x}_i$ are known. The setting presented in Table 3 can also be extended to calibrated d_i^B weights (i.e. d_i^B adjusted for under-coverage, non-contact or non-response; cf. Särndal and Lundström (2005)) but this requires additional developments in the theory about the consistency of the MI, IPW and DR estimators. In the next sections we briefly present the methods implemented in the package.

2.2. Prediction-based approach

Prediction estimators

In the prediction approach the following semi-parametric model for the finite population is assumed:

$$E_\xi \left(y_i \mid \boldsymbol{x}_i \right) = m \left(\boldsymbol{x}_i, \boldsymbol{\beta} \right), \text{ and } \quad V_\xi \left(y_i \mid \boldsymbol{x}_i \right) = v \left(\boldsymbol{x}_i \right) \sigma^2, \quad i = 1, 2, \dots, N, \tag{1}$$

where the mean function $m(\cdot, \cdot)$ and the variance $v(\cdot)$ have known forms and y_i are also assumed to be conditionally independent given the $\boldsymbol{x}_i$. The model in (1) is assumed to hold for all units in the non-probability sample S_A. The parameters of the model in (1) can be estimated using the quasi maximum likelihood estimation method, which includes linear and non-linear models, such as generalized linear models (GLM). Let $\boldsymbol{\beta}_0$ and σ_0^2 be the true values of the model parameters $\boldsymbol{\beta}$ and σ^2 under the adopted model and $\hat{\boldsymbol{\beta}}$ be the quasi maximum likelihood estimator of $\boldsymbol{\beta}_0$. Let $m_i = m(\boldsymbol{x}_i, \boldsymbol{\beta}_0)$ and $\hat{m}_i = m(\boldsymbol{x}_i, \hat{\boldsymbol{\beta}})$ be calculated for all units $i = 1, ..., N$. Under this setting, as Wu (2022) notes, there are two commonly used prediction estimators:

$$\hat{\mu}_{y,PR1} = \frac{1}{N} \sum_{i=1}^{N} \hat{m}_i \quad \text{and} \quad \hat{\mu}_{y,PR2} = \frac{1}{N} \left\{ \sum_{i \in S_A} y_i - \sum_{i \in S_A} \hat{m}_i + \sum_{i=1}^{N} \hat{m}_i \right\}. \tag{2}$$

Under linear models, where $m(\boldsymbol{x}_i, \boldsymbol{\beta}) = \boldsymbol{x}_i'\boldsymbol{\beta}$, the two estimators (2) reduce to:

$$\hat{\mu}_{y,PR1} = \boldsymbol{\mu}_{\boldsymbol{x}}'\hat{\boldsymbol{\beta}} \quad \text{and} \quad \hat{\mu}_{y,PR2} = \frac{n_A}{N} \left(\bar{y}_A - \overline{\boldsymbol{x}}_A'\hat{\boldsymbol{\beta}} \right) + \boldsymbol{\mu}_{\boldsymbol{x}}'\hat{\boldsymbol{\beta}}, \tag{3}$$

where $\boldsymbol{\mu}_{\boldsymbol{x}} = N^{-1} \sum_{i=1}^{N} \boldsymbol{x}_i$ is the vector of the population means of the $\boldsymbol{x}$ variables and $\overline{\boldsymbol{x}}_A = n_A^{-1} \sum_{i \in S_A} \boldsymbol{x}_i$ is the vector of the simple means of $\boldsymbol{x}$ from the non-probability S_A sample. If the linear model contains an intercept and $\hat{\boldsymbol{\beta}}$ is the ordinary least square estimator, then $\hat{\mu}_{y,PR1} = \hat{\mu}_{y,PR2}$.

This form is appealing as it only requires the non-probability sample S_A and reference population means (or totals and population size N). If the population means are unknown, they can be replaced by estimates provided by the reference probability sample S_B, i.e. $\sum_{i=1}^{N} \hat{m}_i$ is replaced with $\sum_{i \in S_B} d_i^B \hat{m}_i$ for (2) and $\boldsymbol{\mu}_{\boldsymbol{x}}$ is replaced by $\hat{\boldsymbol{\mu}}_{\boldsymbol{x}} = \hat{N}_B^{-1} \sum_{i \in S_B} d_i^B \boldsymbol{x}_i$ for (3) where $\hat{N}_B = \sum_{i \in S_B} d_i^B$.

Mass imputation estimators

Model-based prediction estimators of μ can be treated as *mass imputation estimators*, since the information on y_i is missing entirely in the reference probability sample S_B (but $\boldsymbol{x}_i$ is available) and y_i can be imputed based on the non-probability sample as values of $\{(y_i, \boldsymbol{x}_i), i \in S_A\}$ are known. The general form of the MI estimator is given by:

$$\hat{\mu}_{y,MI} = \frac{1}{\hat{N}_B} \sum_{i \in S_B} d_i^B y_i^*, \tag{4}$$

where y_i^* is the imputed value of y_i and $\hat{N}_B$ is defined as previously. Under deterministic regression imputation, the $\hat{\mu}_{y,MI}$ estimator reduces to the estimators in (2).

There are several ways of imputing y_i^* and in the package we have implemented the following MI estimators: the semi-parametric approach based on generalized linear models (MI-GLM),

kernel smoothing (MI-NPAR), nearest neighbour matching (MI-NN) and predictive mean matching (MI-PMM).

The properties of the MI-GLM estimator, where y_i^* are imputed with $\hat{m}_i$ from the semi-parametric model, were studied by Kim, Park, Chen, and Wu (2021). In the **nonprobsvy** package, we account for the following GLM families: `gaussian`, `binomial` and `poisson`. The MI-NPAR estimator was proposed by Chen *et al.* (2022) who studied its finite population properties. Currently, we implement this approach using local polynomial regression using `stats::loess` method. In the next releases we will allow users to use the **np** (Hayfield and Racine 2008) or **KernSmooth** (Wand 2025) packages.

The MI-NN estimator was initially proposed by Rivers (2007) under the name *sample matching* and theoretical properties of the MI-NN estimator for large non-probability samples (big data, i.e. covering a significant part of the target population) were studied by Yang *et al.* (2021). The basic idea of NN matching is as follows: 1) for each i unit in the probability sample S_B find a donor j (or donors) in the S_A sample based on some distance between $\boldsymbol{x}_i$ and $\boldsymbol{x}_j$; 2) use the matched values y_j from S_A to impute missing y_i values in the probability sample S_B. Imputed values of y_i^* depend on the number of selected k neighbours: for $k = 1$, the closest one is selected, and for $k > 1$, one can calculate a simple average over a vector of selected y values. A detailed description of the procedure is presented in Algorithm 1 in the Appendix. The MI-NN estimator suffers from the curse of dimensionality, i.e. asymptotic bias of the MI estimator increases as the number of covariates $\boldsymbol{x}$ increases with a fixed k (Abadie and Imbens 2006; Yang and Kim 2020). The PMM approach has been proposed to overcome this issue.

In the PMM approach, matching is done using predicted values of $\hat{m}_i = m\left(\boldsymbol{x}_i, \hat{\boldsymbol{\beta}}\right)$ instead of $\boldsymbol{x}_i$, thus the NN algorithm is modified as follows: 1) fit the model $m\left(\boldsymbol{x}_i, \boldsymbol{\beta}\right)$ to non-probability S_A sample, 2) assign predicted values $\hat{m}_i$ to all units in S_A and S_B; 3) match all units from the S_B sample to donor units from the S_A sample based on $\hat{m}$ values. The MI-PMM estimator is the same as in the NN approach. Chlebicki, Łukasz Chrostowski, and Beręsewicz (2024) studied properties of two variants of the MI-PMM estimator for non-probability samples: matching predicted to predicted ($\hat{m} - \hat{m}$ matching; denoted as MI-PMM A) and matching predicted to observed ($\hat{m} - y$ matching; denoted as MI-PMM B). Details of the procedure can be found in Algorithm 2 and 3 in the Appendix. Chlebicki *et al.* (2024) also prove the consistency of the MI-PMM A estimator under model mis-specification, i.e. the assumed model may differ from the true one.

Variance estimators for the prediction approach

Variance of the MI estimators can be estimated analytically or using the bootstrap approach. The analytical estimator of the variance of the MI-GLM estimator proposed by Kim *et al.* (2021, p. 950) contains two components: $\hat{V}_1$ (based on the information from both samples S_A and S_B) and $\hat{V}_2$ (based exclusively on the probability sample S_B); for the MI-NN estimator Yang *et al.* (2021) proposed a variance estimator for large S_A samples, which reduces to the part for the probability sample S_B (i.e. the design-based variance estimator of the mean, which can easily be obtained from the **survey** package) and a version for smaller samples can be found in Chlebicki *et al.* (2024); with respect to the variance of MI-PMM estimators, Chlebicki *et al.* (2024) propose formulas which are the same as those for the MI-NN estimators. The analytical variance estimator for the MI-NPAR was proposed by Chen *et al.* (2022).

In the bootstrap approach each bootstrap replication $b = 1, ..., B$ consists of the following steps.

1. Independently:

   - draw a simple random sampling with replacement from the non-probability sample S_A (using `base::sample`),
   - draw a sample according to the declared sampling design from the probability sample S_B (e.g. one can use the `as.svrepdesign` function from the **survey** package).

2. Estimate $\mu_{y,MI}^b$ using an appropriate approach (e.g. MI-GLM, MI-NN or MI-PMM).

After obtaining B bootstrap replicates, estimate variance using the following equation:

$$\hat{V}_{\text{boot}} = \frac{1}{B-1} \sum_{b=1}^{B} \left(\hat{\mu}_y^b - \hat{\mu}_y \right)^2, \tag{5}$$

where $\hat{\mu}_y$ is the mean estimated using either the MI-GLM, MI-NN or MI-PMM estimator.

The above approaches are applied when unit-level data from the probability sample S_B are available. If this is not the case and only population means (or totals and population size) are available, we can estimate the variance of the $\mu_{y,MI-GLM}$ estimator using the first component $\hat{V}_1$ of the Kim *et al.* (2021) variance estimator (replaced by the survey-based population quantities, if available). To estimate the variance of the MI-NN and MI-PMM estimators we only allow the bootstrap approach with known population means. Note that the current version of the **nonprobsvy** does not support the use of replicated weights in the probability sample S_B for any of the estimators discussed in this paper.

2.3. Inverse Probability Weighting

Inverse probability weighting (IPW), another popular estimation approach, involves estimating PS given by $\pi_i^A = P(i \in S_A)$. As in the case of the prediction-based approach, there are two variants of the IPW estimator, given by:

$$\hat{\mu}_{y,IPW1} = \frac{1}{N} \sum_{i \in S_A} \frac{y_i}{\hat{\pi}_i^A} \quad \text{and} \quad \hat{\mu}_{y,IPW2} = \frac{1}{\hat{N}^A} \sum_{i \in S_A} \frac{y_i}{\hat{\pi}_i^A}, \tag{6}$$

where the $\hat{\mu}_{y,IPW1}$ is an adjustment of the Horvitz-Thompson estimator, and the $\hat{\mu}_{y,IPW2}$ is an adjustment of the Hájek estimator, where the estimated population size is given by $\hat{N}^A = \sum_{i \in S_A} (\pi_i^A)^{-1}$. The use of this estimator with respect to non-probability samples is discussed by Lee (2006) and Biffignandi and Bethlehem (2021, chapter 13) and there are several approaches to using propensity scores along with alternative versions of the weights (cf. Elliott and Valliant 2017, section 3). In a article by Chen *et al.* (2020) dedicated to the properties of the estimators in (6), the authors proved their consistency and derived their closed form versions. Wu (2022, section 4.2) argues that the $\hat{\mu}_{y,IPW2}$ estimator performs better than $\hat{\mu}_{y,IPW1}$ even if the population size is known.

The construction of the IPW estimator involves two steps: 1) estimating the PS; and 2) deriving d_i^A, which, in our case, are equal to $1/\pi_i^A$. To estimate the propensity scores $\pi_i^A =$

$\pi(\boldsymbol{x}_i, \boldsymbol{\gamma})$ one can use the likelihood approach assuming that the information about $\boldsymbol{x}_i$ is available for each unit in the population given by (7).

$$\ell(\boldsymbol{\gamma}) = \log \left\{ \prod_{i=1}^{N} \left(\pi_i^A\right)^{R_i} \left(1 - \pi_i^A\right)^{1-R_i} \right\} = \sum_{i \in S_A} \log \left\{ \frac{\pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)}{1 - \pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)} \right\} + \sum_{i=1}^{N} \log\left\{1 - \pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)\right\}. \tag{7}$$

In practice, a function of this form cannot be used because not all units from the population are observed. A more realistic approach consists in using a reference probability sample S_B, which means that the second component of the equation in (7) is replaced, yielding a pseudo log-likelihood function given by (8)

$$\ell^*(\boldsymbol{\gamma}) = \sum_{i \in S_A} \log \left\{ \frac{\pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)}{1 - \pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)} \right\} + \sum_{i \in S_B} d_i^B \log\left\{1 - \pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)\right\}. \tag{8}$$

The maximum pseudo-likelihood estimator $\hat{\boldsymbol{\gamma}}$ can be obtained as the solution to the pseudo score equation, which, under the logit function assumed for π_i^A, is given by (9)

$$\boldsymbol{U}(\boldsymbol{\gamma}) = \sum_{i \in S_A} \boldsymbol{x}_i - \sum_{i \in S_B} d_i^B \pi(\boldsymbol{x}_i, \boldsymbol{\gamma}) \boldsymbol{x}_i. \tag{9}$$

In general, pseudo score functions $\boldsymbol{U}(\boldsymbol{\gamma})$ for true values of model parameters $\boldsymbol{\gamma}_0$ are unbiased under the joint *qp* randomization in the sense that $E_{qp}\left\{\boldsymbol{U}\left(\boldsymbol{\gamma}_0\right)\right\} = \boldsymbol{0}$, which implies that the estimator $\hat{\boldsymbol{\alpha}}$ is consistent for $\boldsymbol{\alpha}_0$ (Wu 2022).

The terms in equation (9) can be replaced by general estimation equations. Let $\boldsymbol{h}(\boldsymbol{x}, \boldsymbol{\gamma})$ be a user-specified vector of functions with the same dimension of $\boldsymbol{\gamma}$ and $\boldsymbol{G}(\boldsymbol{\gamma})$ is defined as

$$\boldsymbol{G}(\boldsymbol{\gamma}) = \sum_{i \in S_A} \boldsymbol{h}\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right) - \sum_{i \in S_B} d_i^B \pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right) \boldsymbol{h}\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right), \tag{10}$$

then solving for $\boldsymbol{G}(\boldsymbol{\gamma}) = \boldsymbol{0}$ with the chosen parametric form of π_i^A and the chosen $\boldsymbol{h}(\boldsymbol{x}, \boldsymbol{\gamma})$ produces the consistent estimator of $\hat{\boldsymbol{\gamma}}$. In the literature, the most commonly considered functions are $\boldsymbol{h}\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right) = \boldsymbol{x}_i$ and $\boldsymbol{h}\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right) = \boldsymbol{x}_i \pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)^{-1}$. Note that if the function $\boldsymbol{h}_i = \boldsymbol{x}_i$, then $\boldsymbol{G}$ reduces to $\boldsymbol{U}$ and for the second variant of the $\boldsymbol{h}$ function we get the following form of function $\boldsymbol{G}$

$$\boldsymbol{G}(\boldsymbol{\theta}) = \sum_{i \in S_A} \frac{\boldsymbol{x}_i}{\pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)} - \sum_{i \in S_B} d_i^B \boldsymbol{x}_i, \tag{11}$$

which can be viewed as *calibrated* IPW and equation (11) only requires the knowledge of population totals for auxiliary variables $\boldsymbol{x}$. Moreover, the use of equation (11) yields a doubly robust estimator under the assumption that the outcome model is linear (Kim and Riddles 2012).

Variance estimators for the inverse probability weighting approach

Chen *et al.* (2020, section 3.2) derived asymptotic variance estimators for both IPW estimators in (6) and presented the plug-in variance estimator for the $\hat{\mu}_{y,IPW2}$ estimator assuming

logistic regression. In the package we have implemented this approach for `logit`, `probit` and `cloglog` link functions. We refer the reader to Wu (2022, section 6.2) and Chrostowski (2024, chapter 3) for more details on how this estimators are derived based on the general estimating equations approach.

Another approach is to use bootstrap, which is essentially the same as the one presented in (5), where $\hat{\mu}_y$ is replaced by one of the estimators in (6).

2.4. Doubly robust approach

The IPW and MI estimators are suited to correctly specified models for PS and outcome regression models, respectively. The DR approach was proposed to improve robustness and efficiency (cf. Robins, Rotnitzky, and Zhao 1994). It incorporates a prediction model for the response variable y_i and a PS model for participation R_i^A. This approach is doubly robust in the sense that the DR estimator remains consistent even if one of the models is mis-specified. We need to consider a joint randomization approach involving a non-probability sample S_A and a probability sample S_B and DR inference is conducted within the qp or ξp framework without specifying which one is correct. The general formula for the for the DR estimator is given by

$$\tilde{\mu}_{y,DR} = \frac{1}{N} \sum_{i \in S_A} \frac{y_i - m_i}{\pi_i^A} + \frac{1}{N} \sum_{i=1}^{N} m_i,$$

where π_i^A and m_i are defined as previously. In the next subsections we discuss two approaches to the DR estimation.

Parameters estimated separately

Chen *et al.* (2020) proposed two DR estimators given in (12) and (13) assuming that the population size is either known or estimated:

$$\hat{\mu}_{y,DR1} = \frac{1}{N} \sum_{i \in S_A} d_i^A \left\{ y_i - m\left(\boldsymbol{x}_i, \hat{\boldsymbol{\beta}}\right) \right\} + \frac{1}{N} \sum_{i \in S_B} d_i^B m\left(\boldsymbol{x}_i, \hat{\boldsymbol{\beta}}\right), \tag{12}$$

and

$$\hat{\mu}_{y,DR2} = \frac{1}{\hat{N}^A} \sum_{i \in S_A} d_i^A \left\{ y_i - m\left(\boldsymbol{x}_i, \hat{\boldsymbol{\beta}}\right) \right\} + \frac{1}{\hat{N}^B} \sum_{i \in S_B} d_i^B m\left(\boldsymbol{x}_i, \hat{\boldsymbol{\beta}}\right), \tag{13}$$

where $d_i^A = \pi\left(\boldsymbol{x}_i, \boldsymbol{\gamma}\right)^{-1}$, $\hat{N}^A = \sum_{i \in S_A} d_i^A$ and $\hat{N}^B = \sum_{i \in S_B} d_i^B$. The estimator in (13), including the estimated population size, has a better performance in terms of bias and the mean squared error and should be used in practice. However, the main limitation is associated with variance estimation, which is discussed at the end of this section.

Chen *et al.* (2020) suggested constructing the estimators in (12) or (13) based on the two models estimated separately, which is an attractive option, since one can specify a different number of variables for the propensity and outcome model. An alternative approach, proposed by Yang, Kim, and Song (2020) and similar to the one described by Kim and Haziza (2014), is discussed in the next subsection.

Minimization of the bias for doubly robust methods

Yang *et al.* (2020) discussed variable selection for a high-dimensional setting and noted that the asymptotic bias of the estimator, which can increase, cannot be controlled. Therefore, according to Yang *et al.* (2020), the idea is to develop equations that can be used to estimate the $\boldsymbol{\beta}$ and $\boldsymbol{\gamma}$ parameters based on the bias of the population mean estimator. In this way the parameters can be estimated in a single step, rather than in two separate steps. First, the authors derived the asymptotic bias of the $\hat{\mu}_{DR}$, assuming $\boldsymbol{h}(\boldsymbol{x}, \boldsymbol{\gamma}) = \boldsymbol{x}\pi(\boldsymbol{x}, \boldsymbol{\gamma})^{-1}$ for the IPW estimator, which is given by equation (14)

$$
\begin{aligned}
\text{bias}\,(\hat{\mu}_{DR}) = E \mid \hat{\mu}_{DR} - \mu \mid &= E\left\{\frac{1}{N}\sum_{i=1}^{N}\left\{\frac{R_i^A}{\pi\,(\boldsymbol{x}_i, \boldsymbol{\gamma})} - 1\right\}\{y_i - m\,(\boldsymbol{x}_i, \boldsymbol{\beta})\}\right\} \\
&+ E\left\{\frac{1}{N}\sum_{i=1}^{N}\left(R_i^B d_i^B - 1\right)m\,(\boldsymbol{x}_i, \boldsymbol{\beta})\right\}.
\end{aligned}
\tag{14}
$$

The goal of this approach is to minimize $\text{bias}\,(\hat{\mu}_{DR})^2$, which consists in solving the following system of empirical equations:

$$
\begin{pmatrix}
\sum_{i=1}^{N} R_i^A \left\{\frac{1}{\pi(\boldsymbol{x}_i, \boldsymbol{\gamma})} - 1\right\}\{y_i - m\,(\boldsymbol{x}_i, \boldsymbol{\beta})\}\,\boldsymbol{x}_i \\
\sum_{i=1}^{N} \frac{R_i^A}{\pi(\boldsymbol{x}_i, \boldsymbol{\gamma})}\dot{m}\,(\boldsymbol{x}_i, \boldsymbol{\beta}) - \sum_{i \in S_B} d_i^B \dot{m}\,(\boldsymbol{x}_i, \boldsymbol{\beta})
\end{pmatrix} = \boldsymbol{0},
\tag{15}
$$

where $\dot{m}\,(\boldsymbol{x}_i, \boldsymbol{\beta}) = \frac{\partial m(\boldsymbol{x}_i, \boldsymbol{\beta})}{\partial \boldsymbol{\beta}}$. The system in (15) can be solved using the Newton–Raphson method. This approach, without variable selection, is equivalent to that proposed by Kim and Haziza (2014) and was extensively discussed by Chen *et al.* (2020) and Wu (2022) in the context of estimating parameters and the variance of the DR estimator. The main limitation of this approach is the possibility that a solution to (15) may not exist unless the two sets of covariates used in the outcome regression model and the PS model have the same dimensions. That is why Yang *et al.* (2020) suggested using this approach on the union of variables from both models (e.g. after variable selection).

In the **nonprobsvy** package we have implemented these approaches not only for $\boldsymbol{h}(\boldsymbol{x}, \boldsymbol{\gamma}) = \boldsymbol{x}_i \pi\,(\boldsymbol{x}_i, \boldsymbol{\gamma})^{-1}$ but also for $\boldsymbol{h}(\boldsymbol{x}, \boldsymbol{\gamma}) = \boldsymbol{x}$ and various link functions for the propensity score model. As noted in the beginning, the choice of either (12) or (13) results in a different approach to estimating variance, which is discussed in the next subsection.

Variance estimators for the doubly robust approach

Yang *et al.* (2020) derived a closed form estimator for (12) but this requires the knowledge of the population and bias correction to obtain a valid estimator for $V_{\xi p}\,(\hat{\mu}_{y,DR} - \mu_y)$ under the outcome regression model ξ. A doubly robust variance estimator for $\hat{\mu}_{y,DR2}$ given by (13) is not yet available in the literature. In the package, to offer the analytical variance estimator of $\hat{\mu}_{y,DR2}$ we simply replace N with estimated $\hat{N}_A$ and $\hat{N}_B$ but we urge caution when using this approach.

Alternatively, one can use the bootstrap approach. Chen *et al.* (2020) demonstrated that the bootstrap approach presented in Section 2.2 performs well in terms of the coverage rate when one of the working models is correctly specified. This is why this approach is recommended for all users.

2.5. Variable selection algorithms

Yang and Kim (2020) point out that it is crucial to use variable selection techniques during estimation, especially when dealing with high-dimensional non-probability samples. Variable selection not only improves model stability and computational feasibility, but also reduces variance, which can increase when irrelevant auxiliary variables are included.

The most popular approaches described in the literature are penalisation methods, such as *Least Absolute Shrinkage and Selection Operator* (LASSO), *Smoothly Clipped Absolute Deviation* (SCAD) or *Minimax Concave Penalty* (MCP), which, thanks to appropriate loss functions, degenerate the coefficients in variables that have no significant effect on the dependent variable (cf. Tibshirani 1996; Breheny and Huang 2011).

The selection procedure for non-probability methods works in a similar way, with loss functions modified to account for external data sources, such as sample or population totals or averages. In particular, the technique consists of two steps: 1) we select the relevant variables using an appropriately constructed loss function (and possibly using the approach shown in (15) to obtain the final estimates of the model parameters); and 2) we construct the selected estimator using variables selected from the first step. For instance, Yang *et al.* (2020) used (16) as a loss function for estimating outcome equation parameters:

$$\text{Loss}\,(\lambda_{\boldsymbol{\beta}}) = \sum_{i=1}^{N} R_i^A \left[y_i - m\left\{\boldsymbol{x}_i, \boldsymbol{\beta}(\lambda_{\boldsymbol{\beta}})\right\}\right]^2, \tag{16}$$

where $m\left\{\boldsymbol{x}_i, \boldsymbol{\beta}(\lambda_{\boldsymbol{\beta}})\right\}$ is the penalised function for the $\boldsymbol{\beta}$ parameters with a tuning parameter $\lambda_{\boldsymbol{\beta}}$ and loss functions for the PS function presented in Table 4, where $\lambda_{\boldsymbol{\gamma}}$ is the tuning parameter.

$\boldsymbol{h}(\boldsymbol{x}_i, \boldsymbol{\gamma})$ function	Loss function $\lambda_{\boldsymbol{\gamma}}$
$\boldsymbol{x}_i$	$\sum_{j=1}^{p} \left(\sum_{i=1}^{N} \left[R_i^A - \frac{R^B \pi\{\boldsymbol{x}_i, \boldsymbol{\gamma}(\lambda_{\boldsymbol{\gamma}})\}}{\pi_i^B} \right] \boldsymbol{x}_{i,j} \right)^2$
$\boldsymbol{x}_i \pi_i(\boldsymbol{x}_i, \boldsymbol{\gamma})^{-1}$	$\sum_{j=1}^{p} \left(\sum_{i=1}^{N} \left[\frac{R_i^A}{\pi\{\boldsymbol{x}_i, \boldsymbol{\gamma}(\lambda_{\boldsymbol{\gamma}})\}} - \frac{R_i^B}{\pi_i^B} \right] \boldsymbol{x}_{i,j} \right)^2$

Table 4: Loss functions for the PS function depending on the $\boldsymbol{h}(\cdot, \cdot)$ function

where R_i^A and R_i^B are indicator functions defining the fact of being included in the non-probability sample S_A and the probability sample S_B, respectively. Yang *et al.* (2020) only discussed the SCAD penalty and the $\boldsymbol{h}(\boldsymbol{x}_i, \boldsymbol{\gamma}) = \boldsymbol{x}_i$ function for the DR estimator. In the **nonprobsvy** package we have extended this approach to the first variant of $\boldsymbol{h}(\boldsymbol{x}_i, \boldsymbol{\gamma})$, shown in the first row of Table (4), and have allowed the user to select other link functions for the π_i^A, implemented other penalty functions and extended the possibility of selecting variables to MI and IPW estimators. In the next section we discuss how to define the approaches presented above.

3. The main function and the package functionalities

3.1. The `nonprob` function

The **nonprobsvy** package is built around the `nonprob` function. The main design objective was to make the use of `nonprob` as similar as possible to standard R functions for fitting statistical models, such as `stats::glm`, while incorporating survey design features from the **survey** package. The most important arguments are given in Table 5 and the obligatory ones include `data` as well as one of the following three –`selection`, `outcome`, or `target` – depending on which method has been selected. In the case of `outcome` and `target` multiple y variables can be specified.

Argument	Description
`data`	a `data.frame` with data from the non-probability sample
`selection`	a `formula` for the selection (PS) equation
`outcome`	a `formula` for the outcome equation (e.g. `y1 + y2`~`x1 + x2`)
`target`	a `formula` with target variables (e.g. `y1+y2+y3`)
`svydesign`	an optional `svydesign2` object
`pop_totals`, `pop_means`, `pop_size`	an optional named `vector` with population totals or means of the covariates and population size
`method_selection`	a link function for the IPW approach (`c("logit", "probit", "cloglog")`)
`method_outcome`	specification of the MI approach (`c("glm", "nn", "pmm", "npar")`)
`family_outcome`	a GLM family for the MI approach (`c("gaussian", "binomial", "poisson")`)
`subset`	an optional `vector` specifying a subset of observations to be used in the fitting process
`strata`	an optional parameter for estimation for sub-populations (currently not supported)
`case_weights`	an optional `vector` of prior case-weights to be used in the fitting process
`na_action`	a `function` indicating what to do with `NA`'s (default `na.omit`)
`control_selection`, `control_outcome`, `control_inference`	control functions with parameters for PS, the outcome model and variance estimation, respectively
`start_selection`, `start_outcome`	an optional `vector` with starting values for the parameters of the PS and the outcome equation
`verbose`	a `logical` value indicating if information should be printed
`se`	a `logical` value indicating whether to calculate and return the standard error of the estimated mean
`...`	Additional optional arguments

Table 5: A description of the `nonprob` function arguments

The package allows the user to provide either reference population data (via the `pop_totals`, or `pop_means` and `pop_size`) or a probability sample declared by the `svydesign` argument (`svydesign2` class from the **survey** package). The `nonprob` function is used to specify inference methods through the `selection` and `outcome` arguments.

If a `svydesign2` object is provided and the `selection` argument is specified, then the IPW estimators are used (by default parameters of the PS model employ (9)), if the `outcome` ar-

gument is specified, then the MI approach is used (the default option is the MI-GLM with the `gaussian` family) and if both are specified, then the DR approach is applied (parameters $(\boldsymbol{\beta}, \boldsymbol{\gamma})$ are estimated separately and the (13) is used). A particular inference method is selected through the `method_selection`, `method_outcome`, `family_outcome`, `control_selection` and `control_outcome` arguments. The variance estimation method is selected through the `control_inference` argument.

Resulting object of class `nonprob` is a list that contains the following (most important) elements:

- `data` – a `data.frame` containing the non-probability sample.

- `X` – a `matrix` containing both samples,

- `y` – a `list` containing all variables declared in either the `target` or `outcome` arguments,

- `R` – a numeric `vector` informing about inclusion in the non-probability sample,

- `ps_scores` – propensity scores or `NULL` (for the MI estimators),

- `ipw_weights` – inverse probability weights or `NULL` (for the MI estimators),

- `output` - a `data.frame` containing point and standard error estimates,

- `confidence_interval` - a `data.frame` containing confidence intervals for the mean,

- `outcome` – a `list` of results for each `outcome` model,

- `selection` – a `list` of results for the `selection` model,

- `svydesign` – a `svydesign2` object passed by the `svydesign` argument.

- `ys_rand_pred, ys_nons_pred` – a `list` of predicted values from the MI model.

3.2. Controlling the type of estimators

The `control_out` function can be used to specify various aspects of the estimation process, including the variable selection methods (through different `penalty` options like SCAD, LASSO, and MCP with their respective tuning parameters defined in the same way as in the `control_sel` function; cf. Yang *et al.* (2020)), and detailed configuration for NN, PMM and NPAR approaches (using parameters like `pmm_match_type`, `pmm_weights`, `pmm_k_choice` or `npar_loess`). For NN and PMM approaches we use the **RANN** package (Jefferis, Kemp, Arya, and Mount 2024) with the kd-tree algorithm and the Euclidean distance as default. We currently do not support other distances (e.g. Gower). Table 6 presents example usage of the `control_out` function for three types of MI estimators.

The `control_sel` function provides essential control parameters for fitting the selection model in the `nonprob` function. It allows users to select between MLE given by (9) or GEE defined in (11) through the `est_method` argument, specify the $\boldsymbol{h}(\cdot, \cdot)$ function through the `gee_h_fun` argument, specify the optimizer (`optimizer` argument) and which variable selection method should be applied (using different penalty functions like SCAD, lasso, and MCP by specifying

Estimator	Declaration with `control_out`

MI-GLM with the LASSO penalty and 5 folds	`nonprob(outcome = y1 ~ x1 + x2,`
	`data = df, svydesign=prob,`
	`control_outcome = control_out(penalty="lasso", folds=5))`

MI-NN with the `bd` algorithm	`nonprob(outcome = y1 ~ x1 + x2,`
	`data = df, svydesign=prob,`
	`control_outcome = control_out(treetype = "bd"))`

MI-PMM A with $k = 3$	`nonprob(outcome = y1 ~ x1 + x2,`
	`data = df, svydesign=prob,`
	`control_outcome = control_out(k=3, pmm_match_type=2))`

Table 6: Example declarations of the MI estimators

the `penalty` argument) along with parameters (e.g. the number of folds through the `nfolds` argument). The parameters of the PS for the calibrated IPW is estimated by using the **nleqslv** package and fitting parameters (arguments starting with the `nleqslv_*`). Table 7 presents example usage of the `control_sel` function for two types of IPW estimators.

Estimator	Declaration with `control_sel`

Calibrated IPW	`nonprob(selection = ~ x1 + x2, target = ~y1,`
	`data = df, svydesign = prob,`
	`control_selection = control_sel(est_method="gee"))`

IPW with the MCP penalty and 5 folds	`nonprob(selection = ~ x1 + x2, target = ~y1,`
	`data = df, svydesign = prob,`
	`control_selection = control_sel(penalty="MCP", nfolds=5))`

Table 7: Example declarations of the IPW estimators

3.3. Controlling variance estimation

Finally, the `control_inf` function configures the parameters for variance estimation in the `nonprob` function. It allows users to specify whether the analytical or bootstrap approach should be used (the `var_method` argument), whether the variable selection method should be applied (the `vars_selection` argument) and what type of bootstrap should be applied for the probability sample (the `rep_type` argument). This function is also used to specify the inference procedure for the DR approach: if a union of variables should be applied (the `vars_combine` argument) and if the bias correction (the `bias_correction` argument) should be applied after variable selection (the `vars_selection`) and variable combination. Table 8 presents example usage of the `control_inf` function for the IPW and DR estimators.

In the next sections we present a case study illustrating the process of integrating a non-probability sample with a reference probability sample. We present various estimators and compare them. Finally, we describe more advanced options available in the package.

Estimator	Declaration with the `control_sel`
Calibrated IPW with variable selection, bootstrap and $B = 50$	`nonprob(selection = ~ x1 + x2, target = ~y1,` `data = df, svydesign = prob,` `control_selection = control_sel(est_method="gee"),` `control_inference = control_inf(vars_selection=TRUE,` `var_method="bootstrap", rep_type="subbootstrap", B=50))`
The DR with the SCAD penalty, 5 folds and bias correction	`nonprob(selection = ~ x1 + x2, outcome = y1 ~ x1 + x2,` `data = df, svydesign = prob,` `control_selection = control_sel(penalty="SCAD", nfolds=5),` `control_inference = control_inf(vars_selection=TRUE,` `bias_correction=TRUE, vars_combine=TRUE))`

Table 8: Example declarations of the IPW estimators

4. Data analysis example

4.1. Description of the data

The package can be installed in the standard manner using:

```
R> install.packages("nonprobsvy")
```

Before we explain the case study let us first load the necessary packages.

```
R> library(nonprobsvy) ## for estimation
R> library(ggplot2)    ## for visualisation
```

The goal of the case study was to integrate survey (`jvs`) and administrative (`admin`) data about job vacancies in Poland. The first source, the Job Vacancy Survey (JVS), contains 6,523 units. The JVS provides a probability sample drawn according to a stratified sampling design. More details can be found in Statistics Poland (2021). The dataset contains information about the NACE code (14 levels, the `nace` column), `region` (16 levels), sector (2 levels, the `private` column), company size (3 levels: Small up to 9, Medium 10-49 and Large 50+) and the final weight (i.e. the design weight corrected for non-contact and non-response), which is treated as the d weight.

```
R> data(jvs)
R> head(jvs)


  id private size nace region weight
1 j_1       0    L    0     14      1
2 j_2       0    L    0     24      6
3 j_3       0    L  R.S     14      1
4 j_4       0    L  R.S     14      1
5 j_5       0    L  R.S     22      1
6 j_6       0    M  R.S     26      1
```

Since the **nonprobsvy** package relies on the functionalities of the **survey** package, we need to define the `svydesign2` object via the `svydesign` function, as shown below. The dataset does not contain the true stratification variable, so we use a simplified version by specifying `~ size + nace + region`; similarly, since we do not have information regarding non-response and its correction, we simply assume that the `weight` sums up to the population size.

```
R> jvs_svy <- svydesign(ids = ~ 1,
+                       weights = ~ weight,
+                       strata = ~ size + nace + region,
+                       data = jvs)
```

Our second source (`admin`), the Central Job Offers Database (CJOD), is a register containing all vacancies submitted to Public Employment Offices (see https://oferty.praca.gov.pl). We treat this register as a *non-probability sample* since it contains administrative data provided on a voluntary basis, so the inclusion mechanism is unknown. This dataset was prepared in such a way that records deemed to be out of scope (either in terms of the definition of vacancy or the population of entities) were excluded. In addition to the same variables found in the JVS, the dataset contains one called `single_shift`, which is our target variable, defined as: **whether a company seeks at least one employee for a single-shift job**. The goal of this case study is to estimate *the share of companies that seek employees for a single-shift job* in Poland in a given quarter.

```
R> data(admin)
R> head(admin)

    id private size nace region single_shift
1 j_1       0    L    P     30        FALSE
2 j_2       0    L    O     14         TRUE
3 j_3       0    L    O     04         TRUE
4 j_4       0    L    O     24         TRUE
5 j_5       0    L    O     04         TRUE
6 j_6       1    L    C     28        FALSE
```

One should keep in mind that this paper does not aim to provide a complete tutorial on how to use non-probability samples for statistical inference. We therefore do not include the stage of aligning variables to meet the same definitions, assessing the strength of the relation between auxiliary variables and the target variable, the selection mechanism and the distribution mis-matches between both samples. In the examples below we assume that there is no overlap between both sources and the naive, reference estimate, given by the mean of the `single_shift` column of `admin`, which is equal to 66.1%.

4.2. Estimation

Propensity score approach

First, we start with the IPW approach, which offers the choice between two estimation methods: MLE (default) and GEE (calibrated to the estimated survey totals). We start by

calling the `nonprob` function, where we define the `selection` argument indicating which variables are to be included, the `target` argument, which specifies the variable of interest, i.e. `single_shift`. The remaining arguments specify the `svydesign` object, the dataset and the link function (`method_selection`).

```
R> ipw_est1 <- nonprob(
+   selection = ~ region + private + nace + size,
+   target = ~ single_shift,
+   svydesign = jvs_svy,
+   data = admin,
+   method_selection = "logit" ## this is the default
+ )
```

In order to get the basic information about the estimated target quantity we can use the `print` method to display the object. It provides information about the methods, source of auxiliary variable, variable selection, the naive (uncorrected estimator) and the corrected along with the standard error and confidence interval. The the last line, i.e. `selected estimator`, is limited to the maximum of 5 variables.

```
R> ipw_est1


A nonprob object
 - estimator type: inverse probability weighting
 - method: logit (mle)
 - auxiliary variables source: survey
 - vars selection: false
 - variance estimator: analytic
 - population size fixed: false
 - naive (uncorrected) estimator: 0.6605
 - selected estimator: 0.7224 (se=0.0421, ci=(0.6399, 0.8048))
```

If we want to see detailed information about the approach, we can use the `summary` method. This function provides information on the sample sizes, sum of the IPW weights along with the distribution of IPW weights (for non-probability sample) and propensity scores (denoted as IPW probabilities) for both samples. A more advanced user may be interested in inspecting the details of the fit by accessing the `selection` element (i.e. `ipw_est1$selection`).

```
R> summary(ipw_est1)


A nonprob_summary object
 - call: nonprob(data = admin, selection = ~region + private + nace +
   size, target = ~single_shift, svydesign = jvs_svy, method_selection = "logit")
 - estimator type: inverse probability weighting
 - nonprob sample size: 9344 (18%)
 - prob sample size: 6523 (12.6%)
 - population size: 51870 (fixed: false)
```

```
 - detailed information about models are stored in list element(s): "selection"
----------------------------------------------------------------
 - sum of IPW weights: 52898.13
 - distribution of IPW weights (nonprob sample):
   - min: 1.1693; mean: 5.6612; median: 4.3334; max: 49.9504
 - distribution of IPW probabilities (nonprob sample):
   - min: 0.0189; mean: 0.2894; median: 0.2525; max: 0.8552
 - distribution of IPW probabilities (prob sample):
   - min: 0.0200; mean: 0.2687; median: 0.2291; max: 0.8552
----------------------------------------------------------------
```

If we want to use the calibrated IPW approach, it is necessary to define the `control_sel` function in the `control_selection` argument by setting the `est_method` argument equal to `gee` (the default is `mle`) and the value of `gee_h_fun`.

```
R> ipw_est2 <- nonprob(
+    selection = ~ region + private + nace + size,
+    target = ~ single_shift,
+    svydesign = jvs_svy,
+    data = admin,
+    method_selection = "logit",
+    control_selection = control_sel(gee_h_fun = 1, est_method = "gee")
+ )
```

Results are comparable to the standard IPW point estimate (70.4 vs 72.2) while the standard error is slightly higher.

```
R> ipw_est2

A nonprob object
 - estimator type: inverse probability weighting
 - method: logit (gee)
 - auxiliary variables source: survey
 - vars selection: false
 - variance estimator: analytic
 - population size fixed: false
 - naive (uncorrected) estimator: 0.6605
 - selected estimator: 0.7042 (se=0.0398, ci=(0.6262, 0.7822))
```

The calibrated IPW significantly improves the balance, which can be accessed via the `check_balance` function:

```
R> data.frame(ipw_mle=check_balance(~size, ipw_est1, 1)$balance,
+             ipw_gee=check_balance(~size, ipw_est2, 1)$balance)

      ipw_mle ipw_gee
sizeL  -367.6       0
sizeM  -228.5       0
sizeS  1624.2       0
```

Notice that neither in the package nor in this paper do we provide a detailed description of the post-hoc results, such as the covariate balance. This can be done using existing CRAN packages, e.g. through the `bal.tab` function from the **cobalt** package (Greifer 2024).

Prediction-based approach

If the user is interested in the prediction-based approach, in particular involving MI estimators, then, they should specify the argument `outcome` as a formula (as in the case of the `glm` function). We allow a single outcome (specified as `y ~ x1 + x2 + ... + xk`) and multiple outcomes (as `y1 + y2 + y3 ~ x1 + x2 + ... + xk`). Note that if the `outcome` argument is specified, then there is no need to specify the `target` argument. By default, the GLM type of an MI estimator is used (i.e. `method_outcome="glm"`). In the code below we show how this type of an MI estimator can be declared.

```
R> mi_est1 <- nonprob(
+   outcome = single_shift ~ region + private + nace + size,
+   svydesign = jvs_svy,
+   data = admin,
+   method_outcome = "glm",
+   family_outcome = "binomial"
+ )
R>
R> mi_est1

A nonprob object
 - estimator type: mass imputation
 - method: glm (binomial)
 - auxiliary variables source: survey
 - vars selection: false
 - variance estimator: analytic
 - population size fixed: false
 - naive (uncorrected) estimator: 0.6605
 - selected estimator: 0.7032 (se=0.0112, ci=(0.6812, 0.7252))
```

In order to employ an MI estimator based on NN matching, one can specify `method_outcome = "nn"` for the nearest neighbours search using all variables specified in the `outcome` argument, `method_outcome = "pmm"` to use PMM or `method_outcome = "npar"` to use non-parametric approach. For the NN and PMM estimators, we employ $k = 5$ nearest neighbours (i.e. `control_out(k=5)`) but the variance of the PMM estimator can be only estimated using bootstrap approach suggested by Chlebicki *et al.* (2024). In the case of the MI-NN estimator there is no need to specify the `family_outcome` argument as no model is estimated underneath.

```
R> set.seed(2024)
R>
R> mi_est2 <- nonprob(
+   outcome = single_shift ~ region + private + nace + size,
```

```
+   svydesign = jvs_svy,
+   data = admin,
+   method_outcome = "nn",
+   control_outcome = control_out(k=5)
+ )
R>
R> mi_est3 <- nonprob(
+   outcome = single_shift ~ region + private + nace + size,
+   svydesign = jvs_svy,
+   data = admin,
+   method_outcome = "pmm",
+   family_outcome = "binomial",
+   control_outcome = control_out(k=5),
+   control_inference = control_inf(var_method = "bootstrap", num_boot = 50)
+ )

Bootstrap variance only for the `pmm` method, analytical version during implementation.
```

Results of both estimators are more or less similar, but it should be noted that the NN version suffers from the curse of dimensionality, so the PMM version seems to be more reliable.

```
R> rbind("NN"=mi_est2$output, "PMM"=mi_est3$output)

        mean          SE
NN  0.6799537 0.01568503
PMM 0.7337228 0.02231178
```

As discussed in Section 2, IPW and MI estimators are asymptotically unbiased only when the model and auxiliary variables are correctly specified. To overcome this problem, the user can turn to doubly robust estimators.

The doubly robust approach

In order to choose doubly robust estimation the user needs to specify both the `selection` and `outcome` arguments. These formulas can be specified with the same or varying number of auxiliary variables. As in the MI approach, we also allow multiple outcomes. In the following example code we have specified the non-calibrated IPW and the MI-GLM estimator.

```
R> dr_est1 <- nonprob(
+   selection = ~ region + private + nace + size,
+   outcome = single_shift ~ region + private + nace + size,
+   svydesign = jvs_svy,
+   data = admin,
+   method_selection = "logit",
+   method_outcome = "glm",
+   family_outcome = "binomial"
+ )
R> dr_est1
```

```
A nonprob object
 - estimator type: doubly robust
 - method: glm (binomial)
 - auxiliary variables source: survey
 - vars selection: false
 - variance estimator: analytic
 - population size fixed: false
 - naive (uncorrected) estimator: 0.6605
 - selected estimator: 0.7035 (se=0.0117, ci=(0.6806, 0.7263))
```

Detailed results can be displayed by using the `summary` function, which prints both information about the distribution of the IPW weights as well as the outcome residuals and predictions. For more details one can access the `"outcome"` and `"selection"` fields in the `nonprob` object.

```
R> summary(dr_est1)
```

```
A nonprob_summary object
 - call: nonprob(data = admin, selection = ~region + private + nace +
    size, outcome = single_shift ~ region + private + nace +
    size, svydesign = jvs_svy, method_selection = "logit", method_outcome = "glm",
    family_outcome = "binomial")
 - estimator type: doubly robust
 - nonprob sample size: 9344 (18%)
 - prob sample size: 6523 (12.6%)
 - population size: 51870 (fixed: false)
 - detailed information about models are stored in list element(s): "outcome" and "selecti
----------------------------------------------------------------
 - sum of IPW weights: 52898.13
 - distribution of IPW weights (nonprob sample):
   - min: 1.1693; mean: 5.6612; median: 4.3334; max: 49.9504
 - distribution of IPW probabilities (nonprob sample):
   - min: 0.0189; mean: 0.2894; median: 0.2525; max: 0.8552
 - distribution of IPW probabilities (prob sample):
   - min: 0.0200; mean: 0.2687; median: 0.2291; max: 0.8552
----------------------------------------------------------------
 - distribution of outcome residuals:
   - single_shift: min: -0.9730; mean: 0.0000; median: 0.1075; max: 0.8564
 - distribution of outcome predictions (nonprob sample):
   - single_shift: min: 0.1436; mean: 0.6605; median: 0.6739; max: 0.9758
 - distribution of outcome predictions (prob sample):
   - single_shift: min: 0.1436; mean: 0.5930; median: 0.5938; max: 0.9785
----------------------------------------------------------------
```

Finally, we can use the bias minimisation approach, as proposed by Yang *et al.* (2020), by specifying the `control_inference = control_inf(bias_correction = TRUE)` argument together with the `vars_combine=TRUE` and `vars_selection=TRUE` as this approach requires variable selection followed by variable union from both equations.

```
R> set.seed(2024)
R> dr_est2 <- nonprob(
+    selection = ~ region + private + nace + size,
+    outcome = single_shift ~ region + private + nace + size,
+    svydesign = jvs_svy,
+    data = admin,
+    method_selection = "logit",
+    method_outcome = "glm",
+    family_outcome = "binomial",
+    control_inference = control_inf(bias_correction = TRUE,
+                                    vars_combine = TRUE,
+                                    vars_selection = TRUE),
+    control_selection = control_sel(nfolds = 4, nlambda = 10),
+    control_outcome = control_out(nfolds = 4)
+ )
R> dr_est2

A nonprob object
 - estimator type: doubly robust
 - method: glm (binomial)
 - auxiliary variables source: survey
 - vars selection: true
 - variance estimator: analytic
 - population size fixed: false
 - naive (uncorrected) estimator: 0.6605
 - selected estimator: 0.7037 (se=0.0104, ci=(0.6833, 0.7240))
```

4.3. Comparison of estimates

Finally, as there is no single method for non-probability samples, we suggest comparing results in a single table or a plot. Figure 1 presents point estimates along with 95% confidence intervals. The various estimators show interesting patterns compared to the naive estimate (red dashed line). Both IPW estimators are characterised with large standard errors and the point estimates over 70%. The MI estimators demonstrate notably different behaviours: while MI-PMM produces the highest point estimate with the widest confidence interval, MI-NN yields the lowest estimate, close to the naive value. Results for the other estimators – MI-GLM and DR (with and without bias minimization) – are clustered together, with similar point estimates and confidence interval widths, suggesting some consensus in their bias correction. All these methods indicate a population parameter higher than the naive estimate, but their relative consistency, provides a certain degree of confidence in their bias correction capabilities.

```
R> df_s <- rbind(cbind(ipw_est1$output, ipw_est1$confidence_interval),
+                cbind(ipw_est2$output, ipw_est2$confidence_interval),
+                cbind(mi_est1$output, mi_est1$confidence_interval),
+                cbind(mi_est2$output, mi_est2$confidence_interval),
```

```
+                     cbind(mi_est3$output, mi_est3$confidence_interval),
+                     cbind(dr_est1$output, dr_est1$confidence_interval),
+                     cbind(dr_est2$output, dr_est2$confidence_interval))
R> rownames(df_s) <- NULL
R>
R> df_s$est <- c("IPW (MLE)", "IPW (GEE)", "MI (GLM)", "MI (NN)",
+                "MI (PMM)", "DR", "DR (BM)")
R>
R> ggplot(data = df_s,
+         aes(y = est, x = mean, xmin = lower_bound, xmax = upper_bound)) +
+   geom_point() +
+   geom_vline(xintercept = mean(admin$single_shift),
+              linetype = "dotted", color = "red") +
+   geom_errorbar() +
+   labs(x = "Point estimator and confidence interval", y = "Estimators") +
+   theme_bw()
```


Figure 1: Comparison of estimates of the share of job vacancies offered on a single-shift

4.4. Advanced usage

Bootstrap approach for variance estimation

In the package we allow the user to estimate the variance of the mean analytically (by default) or using the bootstrap approach, as described in Section 2.2. We use analytical variance

estimators proposed in the papers referenced in Section 2. The calculation of the standard error can be disabled using `nonprob(se=FALSE)`. The bootstrap approach implemented in the package refers to:

- the non-probability sample – currently only simple random sampling with replacement is available via the `base::sample`,

- the probability sample – all the approaches implemented in the `as.svrepdesign` function of the **survey** package are supported and we refer the reader to the relevant help file.

The bootstrap approach is specified via the `control_inf()` function with `var_method = "bootstrap"`. The bootstrap method for the probability sample is controlled via the `rep_type` argument, which passes the method to the `as.svrepdesign` function. The number of iterations is set in the `num_boot` argument (100 by default). If the samples are large or the estimation method is complicated (e.g. involves variable selection) one can set `verbose=TRUE` to track progress. By default bootstrap results are stored in the `boot_sample` element of the resulting list (to disable this option, `keep_boot` should be set to `FALSE`). The following code is an example of applying the IPW approach with the bootstrap approach specified by the argument `control_inference` of the `nonprob` function.

```
R> set.seed(2024)
R> ipw_est1_boot <- nonprob(
+   selection = ~ region + private + nace + size,
+   target = ~ single_shift,
+   svydesign = jvs_svy,
+   data = admin,
+   method_selection = "logit",
+   control_inference = control_inf(var_method = "bootstrap", num_boot = 50),
+   verbose = FALSE
+ )
```

Next, we compare the estimated standard error of variance estimation for the analytical and the bootstrap approach.

```
R> rbind("IPW analytic variance"=ipw_est1$output,
+        "IPW bootstrap variance"=ipw_est1_boot$output)

                            mean         SE
IPW analytic variance  0.7223628 0.04207711
IPW bootstrap variance 0.7223628 0.04307590
```

The boot samples can be accessed via the `boot_sample` element of the output list of the `nonprob` function. Note that the output is returned as a `matrix` because we allow multiple `target` variables.

```
R> head(ipw_est1_boot$boot_sample, n=3)
```

```
      single_shift
[1,]     0.7406651
[2,]     0.6989083
[3,]     0.7667519
```

Variable selection algorithms

In this section we briefly show how to use variable selection algorithms. In order to indicate that a variable selection algorithm should be used one should specify the `control_inference = control_inf(vars_selection = TRUE)` argument. Then, the user should either leave the default setting or specify the outcome parameters via the `control_out` function or the `control_sel` function. Both functions have the same parameters:

- `penalty` – The penalization function used during variables selection (possible values: `c("SCAD", "lasso", "MCP")`)

- `nlambda` – The number of λ values; by default set to 50 (grid search).

- `lambda_min` – The smallest value for λ, as a fraction of `lambda.max`; 0.001 by default.

- `lambda` – A user specified vector of lambdas (only for the `control_sel` function).

- `nfolds` – The number of folds for cross validation; by default set to 10.

- `a_SCAD, a_MCP` – The tuning parameter of the SCAD and MCP penalty for the selection model; by default set to 3.7 and 3, respectively.

In the case of the MI approach we rely on the **ncvreg** package (Breheny and Huang 2011), which is the only R package that employs the SCAD method. For the IPW and DR approaches, we have developed our own codes in C++ via the **Rcpp** and **RcppArmadillo** packages. In the code below we apply variable selection for the MI-GLM estimator using only 5 folds, 25 possible values of λ parameters and the LASSO penalty.

```
R> set.seed(2024)
R> mi_est1_sel <- nonprob(
+   outcome = single_shift ~ region + private + nace + size,
+   svydesign = jvs_svy,
+   data = admin,
+   method_outcome = "glm",
+   family_outcome = "binomial" ,
+   control_outcome = control_out(nfolds = 5, nlambda = 25, penalty = "lasso"),
+   control_inference = control_inf(vars_selection = TRUE),
+   verbose = TRUE
+ )

Starting CV fold #1
Starting CV fold #2
Starting CV fold #3
Starting CV fold #4
Starting CV fold #5
```

In this case study, the MI-GLM estimator with variable selection yields almost the same results as the approach without it. Point estimates and standard errors differ at the fourth and third digit, respectively.

```
R> rbind("MI without var sel"= mi_est1$output,
+        "MI with var sel"   = mi_est1_sel$output)


                          mean         SE
MI without var sel  0.7031991 0.01120162
MI with var sel     0.7019285 0.01102080
```

The result object of the `cv.ncvreg` class is stored in the `"outcome"` element of the result, and you can access the selected variables using the `coef` generic method as follows.

```
R> round(coef(mi_est1_sel$outcome$single_shift), 4)
```

(Intercept)	region04	region06	region08	region10	region12
0.2817	0.0023	0.3272	0.3195	0.2118	0.1773
region14	region16	region18	region20	region22	region24
0.0142	0.0791	0.0000	0.0000	0.0046	-0.2555
region26	region28	region30	region32	private	naceD.E
0.1332	0.0000	0.0000	0.0000	-0.6087	0.1762
naceF	naceG	naceH	naceI	naceJ	naceK.L
1.9175	-0.4556	-0.5605	-1.0964	0.9216	1.0372
naceM	naceN	naceO	naceP	naceQ	naceR.S
1.0027	-0.1839	1.4748	0.5371	-0.7113	-0.8136
sizeM	sizeS				
0.9971	1.5353				

If a user is interested in viewing the PS estimation results, then access to the `"selection"` element is required. This is a `list` of all the results returned by the MLE or GEE estimation method. For example, the evaluation of coefficients can be done in the following way

```
R> round(ipw_est1$selection$coefficients,4)
```

(Intercept)	region04	region06	region08	region10	region12
-0.6528	0.8378	0.1995	0.1048	-0.1576	-0.6099
region14	region16	region18	region20	region22	region24
-0.8415	0.7639	1.1781	0.2225	-0.0375	-0.4067
region26	region28	region30	region32	private	naceD.E
0.2029	0.5786	-0.6102	0.3274	0.0590	0.7727
naceF	naceG	naceH	naceI	naceJ	naceK.L
-0.3778	-0.3337	-0.6517	0.4118	-1.4264	0.0617
naceM	naceN	naceO	naceP	naceQ	naceR.S
-0.4068	0.8003	-0.6935	1.2510	0.3029	0.2223
sizeM	sizeS				
-0.3641	-1.0292				

In the future versions we plan to align the selection and outcome models to return the same type of elements and objects.

5. Classes and `S3Methods`

The package contains the main class `nonprob` and the supplementary class `nonprob_method` and the related `nonprob_summary` class. All available `S3methods` can be obtained by calling `methods(class="nonprobsvy")`. For instance, the `check_balance` function, already mentioned in the case study, is used to view the balance by checking how the PS weights reproduce known or estimated population totals; the `nobs` function returns the sample size of the probability and non-probability samples.

```
R> nobs(dr_est1)


   prob nonprob
   6523    9344
```

Table 9 presents methods implemented for the `nonprobsy` class. We have intentionally limited the number of implemented methods as the goal of the package is to provide point and interval estimates. If users are interested in assessing the quality of the models or the covariate balance, they should use existing R packages.

Function	Description
`check_balance`	returns a list of totals for IPW and DR estimators;
`pop_size`	returns the population size (fixed or not)
`confint`	returns the confidence interval for the target variable(s)
`nobs`	returns the number of observations for both samples
`summary`	returns a `nonprob_summary` class with additional information
`update`	returns an `updated` object
`weights`	returns IPW weights for the IPW and DR estimator

Table 9: `S3Methods` implemented in the **nonprobsvy**

The confidence interval for the target variable(s) can be estimated using the generic `confint` method.

```
R> confint(dr_est1, level = 0.99)


        target lower_bound upper_bound
1 single_shift   0.6734515   0.7334882
```

If a user is interested in assessing the distribution of the IPW weights, we suggest using the `weights` method.

```
R> summary(weights(dr_est1))
```

```
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  1.169   2.673   4.333   5.661   7.178  49.950
```

In the case of mass imputation estimators it is possible to use special functions `method_glm` that are of class `nonprob_method`. An example is given below

```
R> res_glm <- method_glm(
+    y_nons = admin$single_shift,
+    X_nons = model.matrix(~ region + private + nace + size, admin),
+    X_rand = model.matrix(~ region + private + nace + size, jvs),
+    svydesign = jvs_svy)
R>
R> res_glm

Mass imputation model (GLM approach). Estimated mean: 0.7039 (se: 0.0115)
```

For the IPW we have created the function `method_ps`, which returns requested functions for a specific link. This approach is motivated by the following reasons: 1) we allow different estimation techniques, such as maximum likelihood and general estimation equations; 2) the IPW estimator is also used for the DR estimator; and 3) variance estimators depend on whether the IPW or the DR estimator is applied (with combination with the MI estimator).

```
R> method_ps()

Propensity score model with logit link
```

Details can be found in help page.

```
R> ?method_ps()
```

6. Summary and future work

The **nonprobsvy** package provides a comprehensive R software solution that addresses inference challenges connected with non-probability samples by integrating them with probability samples or known population totals/means. As non-probability data sources, like administrative registers, voluntary online panels, and social media data become increasingly available, statisticians need robust methods to produce reliable population estimates. The package implements *state-of-the-art* approaches including mass imputation, inverse probability weighting, and doubly robust methods, each designed to correct selection bias by leveraging auxiliary data. By providing a unified framework and its integration with the **survey** package, the **nonprobsvy** makes complex statistical methods for non-probability samples more accessible, enabling researchers to produce robust estimates even when working with non-representative data.

There are several avenues for future development of the **nonprobsvy** package. One key priority is to implement model-based calibration and additional methods for estimating propensity

scores and weights. The package currently assumes no overlap between probability and non-probability samples, so accounting for potential overlap (e.g., in big data sources and registers) is another important extension. Additional planned developments include handling non-ignorable sample selection mechanisms, developing a theory for maintaining consistency with calibration weights, and supporting multiple non-probability samples from various sources for the purpose of data integration. Further methodological extensions under consideration include empirical likelihood approaches for doubly/multiply robust estimation, integration of machine learning methods like debiased/double machine learning from causal inference, handling measurement errors in big data variables, and expanding the bootstrap approach beyond simple random sampling with replacement.

The package will also be extended to handle the `svyrep.design` class from the **survey** package and the **svrep** package. These developments will enhance its capabilities for handling complex survey data structures and modern estimation challenges.

7. Acknowledgements

A. List of symbols

Symbol	Description
U	Target population of size N
S_A	Non-probability sample
S_B	Probability sample
N	Population size
n_A	Size of non-probability sample
n_B	Size of probability sample
$\hat{N}^A$	Estimated size based on non-probability sample
$\hat{N}^B$	Estimated size based on probability sample
$\boldsymbol{x}_i$	Vector of auxiliary variables for unit i
y_i	Value of the study/target variable for unit i
y_i^*	Imputed value for unit i in S_B
π_i^A	Propensity score for unit i in non-probability sample
π_i^B	Inclusion probability for unit i in probability sample
d_i^A	Inverse probability weight $(1/\pi_i^A)$ for non-probability sample
d_i^B	Design weight $(1/\pi_i^B)$ for probability sample
R_i^A	Indicator of inclusion into non-probability sample
R_i^B	Indicator of inclusion into probability sample
μ	Population mean of target variable y
$\mu_{\boldsymbol{x}}$	Population means of auxiliary variables $\boldsymbol{x}$
$m(\boldsymbol{x}_i, \boldsymbol{\beta})$	Semiparametric model for outcome variable
$\dot{m}(\boldsymbol{x}_i, \boldsymbol{\beta})$	First derivative of the $m(\boldsymbol{x}_i, \boldsymbol{\beta})$ with respect to $\boldsymbol{\beta}$
$\pi(\boldsymbol{x}_i, \boldsymbol{\gamma})$	Propensity score model for R_i^A
$\boldsymbol{\beta}$	Parameter vector for outcome model
$\boldsymbol{\gamma}$	Parameter vector for propensity score model
$\lambda_{\boldsymbol{\beta}}, \lambda_{\boldsymbol{\gamma}}$	Tuning parameters for penalisation methods
$\hat{\mu}_{\boldsymbol{x}}$	Estimator for the population means of auxiliary variables $\boldsymbol{x}$
$\bar{\boldsymbol{x}}_A$	A vector of the sample means of the auxiliary variables $\boldsymbol{x}$ from S_A
$\hat{\mu}_{PR}$	Prediction estimators
$\hat{\mu}_{MI}$	Mass imputation estimator
$\hat{\mu}_{IPW}$	Inverse probability weighting estimator
$\hat{\mu}_{DR}$	Doubly robust estimator
$\hat{V}_{boot}$	Variance estimator based on the bootstrap

Table 10: List of symbols and their descriptions

B. Algorithms for the MI-NN and MI-PMM estimators

Algorithm 1: Mass imputation using the k nearest neighbour algorithm

---

1: If $k = 1$, then for each $i \in S_B$ match $\hat{\nu}(i)$ such that $\hat{\nu}(i) = \underset{j \in S_A}{\arg\min}\, d\left(\boldsymbol{x}_i, \boldsymbol{x}_j\right)$.

2: If $k > 1$, then

$$\hat{\nu}(i, z) = \underset{j \in S_A \setminus \bigcup_{t=1}^{z-1} \{\hat{\nu}(i,t)\}}{\arg\min}\, d\left(\boldsymbol{x}_i, \boldsymbol{x}_j\right)$$

i.e. $\hat{\nu}(i, z)$ is z-th nearest neighbour from the sample S_A.;

3: For each $i \in S_B$, calculate the imputed value as

$$y_i^* = \frac{1}{k} \sum_{t=1}^{k} y_{\hat{\nu}(i,t)}.$$

Algorithm 2: Mass imputation using predictive mean matching variant: $\hat{y} - \hat{y}$ matching

---

1: Estimate regression model $m(\boldsymbol{x}, \boldsymbol{\beta})$ parameters.;

2: Predict

$$\hat{y}_i = m\left(\boldsymbol{x}_i, \hat{\boldsymbol{\beta}}\right), \hat{y}_j = m\left(\boldsymbol{x}_j, \hat{\boldsymbol{\beta}}\right)$$

for $i \in S_B, j \in S_A$ and assign each $i \in S_B$ to $\hat{\nu}(i)$, where

$$\hat{\nu}(i) = \underset{j \in S_A}{\arg\min}\, d\left(\hat{y}_i, \hat{y}_j\right).$$

3: If $k > 1$, then:

$$\hat{\nu}(i, z) = \underset{j \in S_A \setminus \bigcup_{t=1}^{z-1} \{\hat{\nu}(i,t)\}}{\arg\min}\, d\left(\hat{y}_i, \hat{y}_j\right)$$

e.g., $\hat{\nu}(i, z)$ is z-th nearest neighbour from a sample S_A.;

4: For $i \in S_B$, calculate imputation value as

$$y_i^* = \frac{1}{k} \sum_{t=1}^{k} y_{\hat{\nu}(i,t)}.$$

Algorithm 3: Mass imputation using predictive mean matching variant: $\hat{y} - y$ matching

---

1: Estimate regression model $m(\boldsymbol{x}, \boldsymbol{\beta})$ parameters.;
2: Predict

$$\hat{y}_i = m\left(\boldsymbol{x}_i, \hat{\boldsymbol{\beta}}\right)$$

for $i \in S_B$ and assign each $i \in S_B$ do $\hat{\nu}(i)$, where

$$\hat{\nu}(i) = \underset{j \in S_A}{\arg\min}\, d\left(\hat{y}_i, y_j\right).$$

3: If $k > 1$, then:

$$\hat{\nu}(i, z) = \underset{j \in S_A \setminus \bigcup_{t=1}^{z-1}\{\hat{\nu}(i,t)\}}{\arg\min}\, d\left(\hat{y}_i, y_j\right).$$

4: For each $i \in S_B$ calculate imputation value as

$$y_i^* = \frac{1}{k}\sum_{t=1}^{k} y_{\hat{\nu}(i,t)}.$$

References

Abadie A, Imbens GW (2006). "Large sample properties of matching estimators for average treatment effects." *Econometrica*, **74**(1), 235–267.

Beaumont JF (2020). "Are probability surveys bound to disappear for the production of official statistics." *Survey Methodology*, **46**(1), 1–28.

Beręsewicz M (2017). "A two-step procedure to measure representativeness of internet data sources." *International Statistical Review*, **85**(3), 473–493.

Beręsewicz M, Szymkowiak M, Chlebicki P (2025). "Quantile balancing inverse probability weighting for non-probability samples." *Forthcomming to the Survey Methodology*, **51**, 0–0.

Biffignandi S, Bethlehem J (2021). *Handbook of Web Surveys.* John Wiley & Sons. ISBN 9781119371687. doi:10.1002/9781119371717.

Breheny P, Huang J (2011). "Coordinate descent algorithms for nonconvex penalized regression, with applications to biological feature selection." *Annals of Applied Statistics*, **5**(1), 232–253. doi:10.1214/10-AOAS388. URL https://doi.org/10.1214/10-AOAS388.

Castro Martín L (2024). *INPS: Inference from Non-Probability Samples.* URL https://pypi.org/project/inps/.

Chen J, Valliant R, Elliott M (2018). "Model-assisted calibration of non-probability sample survey data using adaptive LASSO." *Survey Methodology*, **44**(1), 117–144.

Chen S, Yang S, Kim JK (2022). "Nonparametric Mass Imputation for Data Integration." *Journal of Survey Statistics and Methodology*, **10**(1), 1–24. ISSN 2325-0984, 2325-0992. doi:10.1093/jssam/smaa036. URL https://academic.oup.com/jssam/article/10/1/1/5983829.

Chen Y, Li P, Wu C (2020). "Doubly robust inference with nonprobability survey samples." *Journal of the American Statistical Association*, **115**(532), 2011–2021.

Chlebicki P, Łukasz Chrostowski, Beręsewicz M (2024). "Data integration of non-probability and probability samples with predictive mean matching." 2403.13750, URL https://arxiv.org/abs/2403.13750.

Chrostowski L (2024). "Statistical inference with non-probability samples." Master's thesis, Adam Mickiewicz Universith, URL https://github.com/ncn-foreigners/graduation-theses/blob/main/2024-Chrostowski.pdf.

Citro CF (2014). "From multiple modes for surveys to multiple data sources for estimates." *Survey Methodology*, **40**(2), 137–162.

Cobo B, Ferri-García R, Rueda-Sánchez JL, Rueda MdM (2024). "Software review for inference with non-probability surveys." *The Survey Statistician*, **90**, 40–47.

Daas PJ, Puts MJ, Buelens B, Hurk PAvd (2015). "Big data as a source for official statistics." *Journal of Official Statistics*, **31**(2), 249–262.

Diallo MS (2021). "samplics: a Python Package for selecting, weighting and analyzing data from complex sampling designs." *Journal of Open Source Software*, **6**(68), 3376. doi:10.21105/joss.03376. URL https://doi.org/10.21105/joss.03376.

Elliott MR, Valliant R (2017). "Inference for Nonprobability Samples." *Statistical Science*, **32**(2). ISSN 0883-4237. doi:10.1214/16-STS598. URL https://projecteuclid.org/journals/statistical-science/volume-32/issue-2/Inference-for-Nonprobability-Samples/10.1214/16-STS598.full.

Gelman A (1997). "Poststratification into many categories using hierarchical logistic regression." *Survey Methodology*, **23**, 127.

Goodrich B, Gabry J, Ali I, Brilleman S (2024). *rstanarm: Bayesian applied regression modeling via Stan.* R package version 2.32.1, URL https://mc-stan.org/rstanarm/.

Greifer N (2024). *cobalt: Covariate Balance Tables and Plots.* R package version 4.5.5, URL https://CRAN.R-project.org/package=cobalt.

Grow A, Perrotta D, Del Fava E, Cimentada J, Rampazzo F, Gil-Clavel S, Zagheni E, Flores RD, Ventura I, Weber I (2022). "Is Facebook's Advertising Data Accurate Enough for Use in Social Science Research? Insights from a Cross-National Online Survey." *Journal of the Royal Statistical Society Series A: Statistics in Society*, **185**(2), 343–363. ISSN 0964-1998. doi:10.1111/rssa.12948. URL https://doi.org/10.1111/rssa.12948.

Hayfield T, Racine JS (2008). "Nonparametric Econometrics: The np Package." *Journal of Statistical Software*, **27**(5), 1–32. doi:10.18637/jss.v027.i05. URL https://www.jstatsoft.org/index.php/jss/article/view/v027i05.

Jefferis G, Kemp SE, Arya S, Mount D (2024). *RANN: Fast Nearest Neighbour Search (Wraps ANN Library) Using L2 Metric.* R package version 2.6.2, URL `https://CRAN.R-project.org/package=RANN`.

Kim JK, Haziza D (2014). "Doubly robust inference with missing data in survey sampling." *Statistica Sinica*, **24**(1), 375–394.

Kim JK, Morikawa K (2023). "An Empirical Likelihood Approach to Reduce Selection Bias in Voluntary Samples." *Calcutta Statistical Association Bulletin*, **75**(1), 8–27. `doi:10.1177/00080683231186488`.

Kim JK, Park S, Chen Y, Wu C (2021). "Combining Non-Probability and Probability Survey Samples Through Mass Imputation." *Journal of the Royal Statistical Society Series A: Statistics in Society*, **184**(3), 941–963. ISSN 0964-1998, 1467-985X. `doi:10.1111/rssa.12696`. URL `https://academic.oup.com/jrsssa/article/184/3/941/7068406`.

Kim JK, Riddles MK (2012). "Some theory for propensity-score-adjustment estimators in survey sampling." *Survey Methodology*, **38**(2), 157–165.

Lee S (2006). "Propensity score adjustment as a weighting scheme for volunteer panel web surveys." *Journal of official statistics*, **22**(2), 329.

Lee S, Valliant R (2009). "Estimation for volunteer panel web surveys using propensity score adjustment and calibration adjustment." *Sociological Methods & Research*, **37**(3), 319–343.

Marra G, Rodicw R (2023). *GJRM: Generalized Joint Regression Modelling.*

Meng XL (2018). "Statistical Paradises and Paradoxes in Big Data (I): Law of Large Populations, Big Data Paradox, and the 2016 US Presidential Election." *Annals of Applied Statistics*, **12**, 685–726. `doi:10.1214/18-AOAS1161SF`.

Rivers D (2007). "Sampling for web surveys." In *Proceedings of the Survey Research Methods Section, Joint Statistical Meetings*, pp. 1–26. American Statistical Association, Alexandria, VA.

Robins JM, Rotnitzky A, Zhao LP (1994). "Estimation of regression coefficients when some regressors are not always observed." *Journal of the American statistical Association*, **89**(427), 846–866.

Rueda MdM, Ferri-García R, Castro L (2020). "The R package NonProbEst for estimation in non-probability surveys." *The R Journal*, **12**, 406–418. ISSN 2073-4859. Https://rjournal.github.io/.

Sarig T, Galili T, Eilat R (2023). "balance – a Python package for balancing biased data samples." `2307.06024`, URL `https://arxiv.org/abs/2307.06024`.

Särndal CE, Lundström S (2005). *Estimation in surveys with nonresponse.* John Wiley & Sons.

Schonlau M, Couper MP (2017). "Options for Conducting Web Surveys." *Statistical Science*, **32**(2), 279 – 292. `doi:10.1214/16-STS597`. URL `https://doi.org/10.1214/16-STS597`.

Statistics Poland (2021). "The Demand for :abour: Methodological report." *Methodological report*, Statistical Office in Bydgoszcz, Bydgoszcz, Warsaw. URL https://stat.gov.pl/obszary-tematyczne/rynek-pracy/popyt-na-prace/zeszyt-metodologiczny-popyt-na-prace,3,1.html.

Tibshirani R (1996). "Regression shrinkage and selection via the lasso." *Journal of the Royal Statistical Society Series B: Statistical Methodology*, **58**(1), 267–288.

Virtanen P, Gommers R, Oliphant TE, Haberland M, Reddy T, Cournapeau D, Burovski E, Peterson P, Weckesser W, Bright J, van der Walt SJ, Brett M, Wilson J, Millman KJ, Mayorov N, Nelson ARJ, Jones E, Kern R, Larson E, Carey CJ, Polat İ, Feng Y, Moore EW, VanderPlas J, Laxalde D, Perktold J, Cimrman R, Henriksen I, Quintero EA, Harris CR, Archibald AM, Ribeiro AH, Pedregosa F, van Mulbregt P, SciPy 10 Contributors (2020). "SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python." *Nature Methods*, **17**, 261–272. doi:10.1038/s41592-019-0686-2.

Wand M (2025). *KernSmooth: Functions for Kernel Smoothing Supporting Wand and Jones (1995)*. R package version 2.23-26, URL https://CRAN.R-project.org/package=KernSmooth.

Wu C (2022). "Statistical inference with non-probability survey samples." *Survey Methodology*, **48**, 283–311.

Yang S, Kim JK (2020). "Asymptotic theory and inference of predictive mean matching imputation using a superpopulation model framework." *Scandinavian Journal of Statistics*, **47**(3), 839–861. ISSN 0303-6898, 1467-9469. doi:10.1111/sjos.12429. URL https://onlinelibrary.wiley.com/doi/10.1111/sjos.12429.

Yang S, Kim JK, Hwang Y (2021). "Integration of data from probability surveys and big found data for finite population inference using mass imputation." *Survey Methodology*, **47**, 29–58.

Yang S, Kim JK, Song R (2020). "Doubly Robust Inference when Combining Probability and Non-Probability Samples with High Dimensional Data." *Journal of the Royal Statistical Society Series B: Statistical Methodology*, **82**(2), 445–465. ISSN 1369-7412, 1467-9868. doi:10.1111/rssb.12354. URL https://academic.oup.com/jrsssb/article/82/2/445/7056072.

Affiliation:

Łukasz Chrostowski
Analyx
Krysiewicza 2
61-887 Poznań
E-mail: lukaszchrostowski6@gmail.com

Piotr Chlebicki
Stockholm University
Matematiska institutionen
Albano hus 1
106 91 Stockholm, Sweden
E-mail: piotr.chlebicki@math.su.se
URL: https://github.com/Kertoo, https://www.su.se/profiles/pich3772

Maciej Beręsewicz
Poznań University of Economics and Business
Statistical Office in Poznań

Poznań University of Economics and Business
Department of Statistics
Institute of Informatics and Quantitative Economics
Al. Niepodległosci 10
61-875 Poznań, Poland

Centre for the Methodology of Population Studies
Statistical Office in Poznań
ul. Wojska Polskiego 27/29
60-624 Poznań, Poland
E-mail: maciej.beresewicz@ue.poznan.pl
URL: https://github.com/BERENZ, https://ue.poznan.pl/en/people/dr-maciej-beresewicz/